



Lecture 5

Database Objects

(part 1)





- Constraints in Relational Algebra
- Keys and Foreign Keys
- Constraints on Attributes and Tuples
- Modification of Constraints
- Assertions
- Triggers

Relational Algebra as a Constraint Language

- Two ways

- R is an expression of relational algebra

$$R = \emptyset$$

- The value of R must be empty

- There are no tuples in the result of R

- R and S are expressions of relational algebra

$$R \subseteq S$$

- Every tuple in the result of R must also be in the result of S

Referential Integrity Constraints

- Referential integrity constraint 参照完整性

$$\pi_A(R) \subseteq \pi_B(S)$$

- If we have any value v as the component in attribute A of some tuple in one relation R , then because of our design intentions we may expect that v will appear in a particular component, say for attribute B , of some tuple of another relation S

Referential Integrity Constraints (Cont'd)

- Example

Movies (title, year, length, genre, studioName, **producerC#**)

MovieExec (name, address, **cert#**, netWorth)

$$\pi_{\text{producerC\#}}(\text{Movies}) \subseteq \pi_{\text{cert\#}}(\text{MovieExec})$$

- The **producerC#** of each **Movies** tuple must also appear in the **cert#** of some **MovieExec** tuple

Referential Integrity Constraints (Cont'd)

- Example

StarsIn (**movieTitle**, **movieYear**, starName)

Movies (**title**, **year**, length, genre, studioName, producerC#)

$$\pi_{movieTitle, movieYear}(StarsIn) \subseteq \pi_{title, year}(Movies)$$

Key Constraints

- Example

MovieStar (**name**, address, gender, birthdate)

$$\sigma_{MS1.name=MS2.name \text{ AND } MS1.address \neq MS2.address} (MS1 \times MS2) = \emptyset$$

- No two tuples agree on the *name* component
 - If two tuples agree on *name*, then they must also agree on *address*
- MS1 is shorthand for the renaming

$$\rho_{MS1(name,address,gender,birthdate)} (MovieStar)$$

Additional Constraint Examples

- Example

MovieStar (name, address, **gender**, birthdate)

$$\sigma_{gender \neq 'F' \text{ AND } gender \neq 'M'}(MovieStar) = \emptyset$$

- The set of tuples in *MovieStar* whose *gender* is equal to neither 'F' nor 'M' is empty

Additional Constraint Examples (Cont'd)

- Example

MovieExec (name, address, cert#, netWorth)

Studio (name, address, presC#)

$$\sigma_{netWorth < 10000000} (Studio \bowtie_{presC\# = cert\#} MovieExec) = \emptyset$$

- One must have a net worth of at least \$10,000,000 to be the president of a movie studio

$$\pi_{presC\#} (Studio) \subseteq \pi_{cert\#} (\sigma_{netWorth \geq 10000000} (MovieExec))$$



Constraints and Triggers

- A *constraint* is a relationship among data elements that the DBMS is required to enforce.
 - *Example*: key constraints.
- *Triggers* are only executed when a specified condition occurs, e.g., insertion of a tuple.
 - Easier to implement than complex constraints.



Kinds of Constraints

- **Keys**, (Primary key, Unique Key).
 - Uniqueness of value for a particular attribute or attributes
- **Foreign-key**, or referential-integrity.
- **Attribute-based** constraints.
 - Constrain values of a particular attribute.
 - **Not-Null** Constraints
- **Tuple-based** constraints.
 - Relationship among components.
- **Assertions**: any SQL boolean expression.



Review: Single-Attribute Keys

- Place **PRIMARY KEY** or **UNIQUE** after the type in the declaration of the attribute.
- **Example:**

```
CREATE TABLE movieexec (  
    name          CHAR(30),  
    address       VARCHAR(255),  
    cert          INT          PRIMARY KEY,  
    netWorth      INT  
);
```



Review: Multiattribute Key

- The **title** and **year** together are the key for **movies**:

```
CREATE TABLE movies (  
    title          CHAR(100),  
    year           INT,  
    length         INT,  
    genre          CHAR(10),  
    studioName     CHAR(30),  
    producerC     INT,  
    PRIMARY KEY (title, year)  
);
```



Foreign Keys

- Values appearing in attributes of one relation must appear in certain attributes of another relation.
- **Example:** in `studio(name, address, presC)`, we might expect that a `presC` value also appears in `movieexec.cert`



Expressing Foreign Keys

- To declare a Foreign Key, Use keyword **REFERENCES**, either:
 - With Attribute : After an attribute (for one-attribute keys).
REFERENCES <relation> (<attributes>)
 - As an element of the schema:
FOREIGN KEY (<list of attributes>)
REFERENCES <relation> (<attributes>)
- Referenced attributes must be declared **PRIMARY KEY** or **UNIQUE**.



Example: With Attribute

```
CREATE TABLE movieexec (
```

```
    name          CHAR(30),
```

```
    address       VARCHAR(255),
```

```
    cert          INT          PRIMARY KEY,
```

```
    netWorth      INT
```

```
);
```

```
CREATE TABLE studio (
```

```
    name          CHAR(50)    PRIMARY KEY,
```

```
    address       VARCHAR(255),
```

```
    presC         INT          REFERENCES movieexec(cert)
```

```
);
```




Example: As Schema Element

```
CREATE TABLE movieexec (  
    name          CHAR(30) ,  
    address       VARCHAR(255) ,  
    cert          INT          PRIMARY KEY ,  
    netWorth      INT  
);  
  
CREATE TABLE studio (  
    name          CHAR(50)   PRIMARY KEY ,  
    address       VARCHAR(255) ,  
    presC         INT ,  
    FOREIGN KEY(presC) REFERENCES movieexec(cert)  
);
```



Essence behind of Foreign Keys

studio  movieexec

- Values appearing in attributes of `studio.presC` must appear in `movieexec.cert`.

$$[\text{studio.presC}] \subseteq [\text{movieexec.cert}] + \text{NULL}$$

- Modification on **studio** and **movieexec**

	studio	movieexec
Delete	OK!	?
Insert	?	OK!
Update	?	?



Enforcing Foreign-Key Constraints

- If there is a foreign-key constraint from relation **studio (Referencing)** to relation **movieexec (Referenced)**, two violations are possible:
 1. An insert or update to **studio** introduces values not found in **movieexec**.
 2. A deletion or update to **movieexec** causes some tuples of **studio** to “dangle.”



Actions Taken --- (1)

- An insert or update to **studio** that introduces a nonexistent **executive** in **movieexec** must be

Rejected!



Actions Taken --- (2)

A deletion or update to **movieexec** that removes a **presC** value found in some tuples of **studio** can be handled in three ways.

1. *Default* : Reject the modification.
2. *Cascade* : Make the same changes in **studio**.
 - Deleted executive: delete **studio** tuple.
 - Updated executive: change value in **studio**.
3. *Set NULL* : Change the **studio**.presC to **NULL**.



Example: Cascade

- Delete the **cert=123** tuple from **movieexec**:
 - Then delete all tuples from **studio** that have **presC=123**.
- Update the **cert=345** tuple by changing **345** to **346**:
 - Then change all **studio** tuples with **presC=345** to **presC=346**.



Example: Set NULL

- Delete the **cert=123** tuple from **movieexec**:
 - Change all tuples of **studio** that have **presC=123** to have **presC=NULL**.
- Update the **cert=345** tuple by changing **345** to **34**:
 - Change all tuples of **studio** that have **presC=345** to have **presC=NULL**.



Choosing a Policy



- When we declare a foreign key, we may choose policies **SET NULL** or **CASCADE** independently for deletions and updates.
- Follow the foreign-key declaration by:
ON [UPDATE, DELETE]
[SET NULL, CASCADE]
- Two such clauses may be used.
- Otherwise, the default (**reject**) is used.



Example: Setting Policy

```
CREATE TABLE studio (  
    name CHAR(50) PRIMARY KEY,  
    address VARCHAR(255),  
    presC INT  
REFERENCES movieexec(cert)  
ON DELETE SET NULL  
ON UPDATE CASCADE  
);
```



Deferred Checking of Constraints

- The situation:
 - Arnold Schwarzenegger decides to found a movie studio, called La Vista, of which he is the president
 - violate the foreign-key constraint

```
INSERT INTO studio  
VALUES ('La Vista', 'New York', 23456)
```

We are in trouble because there is no tuple of **movieexec** with certificate number 23456



Solution



- 1) First to insert the tuple for 'La Vista' without a president's certificate

```
INSERT INTO studio(name, address)
```

```
VALUES ('La Vista', 'New York') ;
```

Set presC
as Null

- 2) Then insert a tuple for Arnold Schwarzenegger into **movieexec**

Any problem for this solution?

- 3) Update **Studio**

```
UPDATE Studio
```

```
SET presC# = 23456
```

```
WHERE name = 'La Vista';
```

1、**studio.presC** cannot be PK

2、**Studio.presC** cannot have a Not-Null Constraint



Circular Constraints

- There are cases of **circular constraints**
 - That cannot be fixed by judiciously ordering the database modification steps
 - Impossible to insert new studios with new presidents
 - We cannot insert a tuple with a new value of **presC** into **studio**. Violate the foreign-key **presC**→**movieexec(cert)**
 - We cannot insert a tuple with a new value of **cert** into **movieexec**. Violate the foreign-key **cert**→**studio(presC)**



Deferred Checking



- First, group several SQL statements (the two insertions - one into **studio** and the other into **movieexec**) into a single **transaction**
- Then tell the DBMS not to check the constraints until after the whole transaction has finished its actions and is about to commit.

```
CREATE TABLE Studio (  
    name CHAR(30) PRIMARY KEY,  
    address VARCHAR(255),  
    presC# INT  
        REFERENCES MovieExec (cert#)  
    DEFERRABLE INITIALLY DEFERRED );
```



Deferred Checking



- To inform the DBMS about point (2) in the previous slide
 - The declaration of any constraint may be followed by one of **DEFERRABLE** or **NOT DEFERRABLE** (default)
 - **NOT DEFERRABLE** (default)
 - Every time a database modification statement is executed, the constraint is checked immediately afterwards
 - **DEFERRABLE**
 - We have the option of having it wait until a transaction is complete before checking the constraint
 - **DEFERRABLE INITIALLY DEFERRED**
 - Checking will be deferred to just before each transaction commits
 - **DEFERRABLE INITIALLY IMMEDIATE**
 - The check will be made immediately after each statement



Deferred Checking



- Two additional points about deferring constraints
 - Constraints of any type can be given names
 - If a constraint has a name, say **MyConstraint**, then we can change a deferrable constraint from immediate to deferred by the SQL statement

SET CONSTRAINTS MyConstraint DEFERRED;

- And we can reverse the process by replacing **DEFERRED** in the above to **IMMEDIATE**



Deferred Checking



- **NOT DEFERRABLE** (default)

- 不可延迟，并且约束也无法更改为可延迟状态。约束会在每一句sql statement 之后都进行 check，不符合则 roll back

- **DEFERRABLE**

可延迟状态，在 deferrable 状态时又有2个选项

- **DEFERRABLE INITIALLY DEFERRED**

- 约束会在整个事务进行commit 时check，如果不符合则roll back

- **DEFERRABLE INITIALLY IMMEDIATE**

- 即约束会在每一句sql statement 之后都进行 check，效果等同于not deferrable，但是可以修改延迟状态

在deferrable状态下，可以通过set constraints修改 immediate 或者deferred



OUTLINES



7.1 Keys and Foreign Keys



7.2 Constraints on Attributes and Tuples



7.3 Modification of Constraints



Not-Null Constraints

- Constraints on the value of a particular attribute: **NOT NULL**.

- Disallow tuples in which this attribute is **NULL**

```
CREATE TABLE studio (  
    name          CHAR(50) PRIMARY KEY,  
    address       VARCHAR(255),  
    presC         INT          REFERENCES  
                    movieexec(cert) NOT NULL  
);
```



Not-Null Constraints

- **NOT NULL** has two consequences:
 - We could not insert a tuple into Studio by specifying only the name and address

```
INSERT INTO studio(name, address) VALUES ('La Vista', 'New York');
```
 - We could not use the set-null policy which tells the system to fix foreign-key violations by making **presc** be **NULL**



Attribute-Based Checks

- Constraints on the value of a particular attribute.
- Add **CHECK(<condition>)** to the declaration for the attribute.
 - The condition can be anything that could follow WHERE in a SQL query.
- The condition may use the name of the attribute, but **any other relation or attribute name must be in a subquery** (not supported in pg).



Example: Attribute-Based Check

```
CREATE TABLE studio (  
    name          CHAR(50)      PRIMARY KEY,  
    address       VARCHAR(255),  
    presc         INT           REFERENCES movieexec(cert)  
        CHECK (presc >= 100000)  
);  
  
CREATE TABLE moviestar (  
    name          CHAR(30) PRIMARY KEY,  
    address       VARCHAR(255),  
    gender        CHAR(1) CHECK (gender IN ('F', 'M')),  
    birthdate     CHAR(10)  
);
```



Timing of Attribute-Based Check

- Attribute-based checks are performed **only when a value for that attribute is inserted or updated**.
 - **Example**: CHECK (`presc` \geq 100000) checks every new **presc** and rejects the modification (for that tuple) if the price is < 1000000 .



Attribute-Based Check VS. FK

```
CREATE TABLE studio (  
    name          CHAR(50)          PRIMARY KEY,  
    address       VARCHAR(255) ,  
    presc         INT,  
    CHECK (presc IN (SELECT cert FROM movieexec))  
);
```

- 1、Modification `studio.presc` **NOT IN** `movieexec.cert` **Reject!** As the effect of FK
- 2、Modify `studio.presc` to **NULL** (if NO NULL exist in `movieexec.cert`. Impossible! PK) **Reject!** Not as the effect of FK
- 3、Modification `movieexec` (Delete or Update) is **invisible** to the CHECK in `studio`. **Not as the effect of FK.** This will result in the CHECK constraint becoming violated.



Tuple-Based Checks



- CHECK (<condition>) may be added as a relation-schema element.
- The condition may refer to any attribute of the relation.
 - But other attributes or relations require a subquery (not supported in pg).



Example: Tuple-Based Check

- This constraint is true for every female movie star and for every star whose name does not begin with 'Ms.':

```
CREATE TABLE moviestar (  
    name          CHAR(30) PRIMARY KEY,  
    address       VARCHAR(255) ,  
    gender        CHAR(1) ,  
    birthdate     CHAR(10) ,  
    CHECK (gender = 'F' OR name NOT LIKE 'Ms.%')  
);
```



Timing of Tuple-Based Checks

- Checked every time a tuple is inserted into R and every time a tuple of R is updated
 - If false for the tuple (violated), and insertion or update is rejected!
 - Like an attribute-based CHECK, a tuple-based CHECK is invisible to other relations
 - Even a deletion from R can cause the condition to become false, if R is mentioned in a subquery
 - If a tuple-based check does not have subqueries, then we can rely on its always holding



Comparison of Attribute-, Tuple-Based Checks

Type of Constraint	Where Declared	When Activated	Guaranteed to Hold?
Attribute-based CHECK	With attribute	On insertion to relation or attribute update	Not if subqueries
Tuple-based CHECK	Element of relation schema	On insertion to relation or tuple update	Not if subqueries

- If a constraint on a tuple involves **more than one attribute** of that tuple, then it **must be written as a tuple-based constraint**
- If the constraint involves **only one attribute** of the tuple, then it can be written as **either a tuple- or attribute-based constraint**
 - But the tuple-based constraint will be checked **more frequently** than the attribute-based constraint. **Probably less efficient**



Modification of constraints

- Giving name

```
CREATE TABLE moviestar (  
    name      CHAR(30) CONSTRAINT nameiskey PRIMARY KEY,  
    address   VARCHAR(255),  
    gender    CHAR(1) CONSTRAINT noandro  
              CHECK (gender IN ('F', 'M')),  
    birthdate CHAR(10),  
              CONSTRAINT righttitle CHECK (gender = 'F' OR name NOT  
LIKE 'Ms. %')  
);
```



Modification of constraints

- Altering constraints on tables
 - **SET CONSTRAINT** constraintName **DEFERRED** (or **IMMEDIATE**) ;
 - **ALTER TABLE** relationName **DROP CONSTRAINT** constraintName
 - **ALTER TABLE** relationName **ADD CONSTRAINT** constraintName **CHECK (...)**
- Note
 - The added constraint must be of a kind that can be associated with tuples
 - Tuple-based constraints, key, or foreign-key
 - You cannot add a constraint to a table unless it holds at that time for every tuple in the table

Assertions

- These are database-schema elements, like relations or views.
- Defined by:

CREATE ASSERTION <name>

CHECK (<condition>);

- Condition may refer to any relation or attribute in the database schema.
- An assertion is a boolean-valued SQL expression that must be true at all times.
- **DROP ASSERTION** <name>

Example: Assertion

MovieExec (name, address, cert#, networth)

Studio(name, address, presC#)

No one can become the president of a studio
unless their net worth is at least \$10,000,000

```
CREATE ASSERTION RichPres CHECK  
  (NOT EXISTS  
    (SELECT * FROM studio, movieexec  
      WHERE presc = cert AND networth < 10000000)  
  );
```

Timing of Assertion Checks

- In principle, we must check every assertion after every modification to any relation of the database.
- A clever system can observe that only certain changes could cause a given assertion to be violated.

Comparison of constraints

Type of Constraint	Where Declared	When Activated	Guaranteed to Hold?
Attribute-based CHECK	With attribute	On insertion to relation or attribute update	Not if subqueries
Tuple-based CHECK	Element of relation schema	On insertion to relation or tuple update	Not if subqueries
Assertion	Element of database schema	On any change to any mentioned relation	Yes

Triggers

- *Triggers*

- are only awakened when certain **events**, specified by the database programmer, occur
 - Insert, delete, or update to a particular relation
- Once awakened by its triggering event, the trigger tests a **condition**
 - If the condition does not hold, then nothing else associated with the trigger happens
- If the condition of the trigger is satisfied, the **action** associated with the trigger is performed by the DBMS

Event-Condition-Action Rules

- Another name for “trigger” is *ECA rule*, or *event-condition-action* rule.
- *Event*: typically a type of database modification
- *Condition*: Any SQL boolean-valued expression.
- *Action*: Any SQL statements.

Triggers in SQL

- The check of the trigger's condition and the action of the trigger may be executed
 - either on the state of the database that exists before the triggering event is itself executed
 - or on the state that exists after the triggering event is executed

Triggers in SQL

- The condition and action can refer to both old and/or new values of tuples that were updated in the triggering event
- It is possible to define update events that are limited to a particular attribute or set of attributes

Triggers in SQL

- The programmer has an option of specifying that the trigger executes either
 - Once for each modified tuple, or
 - Row-level trigger
 - Once for all the tuples that are changed in one SQL statement
 - Statement-level trigger

Example: Trigger Definition

```
CREATE TRIGGER NetWorthTrigger
  AFTER UPDATE OF netWorth ON MovieExec
  REFERENCING
    OLD ROW AS OldTuple,
    NEW ROW AS NewTuple
  FOR EACH ROW
  WHEN (OldTuple.netWorth >
        NewTuple.netWorth)
  UPDATE MovieExec
  SET netWorth = OldTuple.netWorth
  WHERE cert = NewTuple.cert;
```

Triggers in PostgreSQL

- The trigger can be specified to fire
 - before the operation is attempted on a row
 - before constraints are checked and the INSERT, UPDATE, or DELETE is attempted
 - or after the operation has completed
 - after constraints are checked and the INSERT, UPDATE, or DELETE has completed
 - instead of the operation
 - in the case of inserts, updates or deletes on a view

Triggers in PostgreSQL

- If the trigger fires before or instead of the event
 - the trigger can skip the operation for the current row
 - or change the row being inserted
- If the trigger fires after the event
 - all changes, including the effects of other triggers, are "visible" to the trigger.

Triggers in PostgreSQL

- A trigger that is marked
 - **FOR EACH ROW** is called once for every row that the operation modifies.
 - **FOR EACH STATEMENT** only executes once for any given operation, regardless of how many rows it modifies
 - an operation that modifies zero rows will still result in the execution of any applicable FOR EACH STATEMENT triggers

Triggers in PostgreSQL

- Triggers that are specified to fire **INSTEAD OF** the trigger event
 - must be marked **FOR EACH ROW**
 - and can only be defined on views
- **BEFORE** and **AFTER** triggers on a view
 - must be marked as **FOR EACH STATEMENT**

Triggers in PostgreSQL

- A trigger definition can specify a Boolean **WHEN** condition
 - which will be tested to see whether the trigger should be fired
 - In row-level triggers the WHEN condition can examine the old and/or new values of columns of the row
 - Statement-level triggers can also have WHEN conditions, although the feature is not so useful for them since the condition cannot refer to any values in the table

Triggers in PostgreSQL

- In FOR EACH ROW triggers
 - the **WHEN** condition can refer to columns of the old and/or new row values by writing **OLD**.column_name or **NEW**.column_name respectively
 - **INSERT** triggers cannot refer to **OLD**
 - **DELETE** triggers cannot refer to **NEW**
 - **INSTEAD OF** triggers do not support **WHEN** conditions

Example : Suppose we want to prevent the average net worth of movie executives from dropping below \$500,000. This constraint could be violated by an insertion, a deletion, or an update to the networth column of

MovieExec(name, address, cert#, networth)

```
CREATE TRIGGER AvgNetWorthTrigger
AFTER UPDATE OF networth ON MovieExec
REFERENCING
    OLD TABLE AS OldStuff,
    NEW TABLE AS NewStuff
FOR EACH STATEMENT
WHEN (500000 > (SELECT AVG (networth) FROM MovieExec))
BEGIN
    DELETE FROM MovieExec
        WHERE (name, address, cert#, networth) IN Newstuff;
    INSERT INTO MovieExec
        (SELECT * FROM Oldstuff);
END;
```

Example of **BEFORE** Trigger

Movies(title, year, length, genre, studioName, producerC#)

```
CREATE TRIGGER FixYearTrigger
BEFORE INSERT ON Movies
REFERENCING
    NEW ROW AS NewRow,
    NEW TABLE AS NewStuff
FOR EACH ROW
WHEN NewRow.year IS NULL
UPDATE NewStuff SET year=1915;
```

Summary

- Referential-integrity constraints
- Attribute-based check constraints
- Tuple-based check constraints
- Modifying constraints
- Assertions
- Invoking the checks
- Triggers