



Lecture 5

Database Objects

(part 2)



NIKON D90 F3.5 3s ISO320



C5000Z F2.8 1/800s ISO50





OUTLINES



8.1

Virtual Views

8.2

Modifying Views



OUTLINES



8.1

Virtual Views

8.2

Modifying Views



Views



- A *view* is a relation defined by a query over stored tables (called *base tables*) and other views.
- Two kinds:
 1. *Virtual* = not stored in the database; just a query for constructing the relation.
 2. *Materialized* = actually constructed and stored.



Declaring Views

- Declared by:

```
CREATE VIEW <view-name>  
AS <SQL query>;
```

- Notice: Relations(Base relations),
Tables(Base tables),
Views

Their differences

Example : Suppose we want to have a view that is a part of the
Movie(title, year , length , incolor, studioName, producerC)
Relation, specifically, the titles and years of the movies made by Paramount
Studios.

```
CREATE VIEW ParamountMovies AS
    SELECT title, year
    FROM Movies
    WHERE studioName = 'Paramount';
```



Example: Our goal is a relation MovieProd with movie titles and the names of their producers.

```
CREATE VIEW MovieProd AS
    SELECT title, name
    FROM Movies, MovieExec
    WHERE producerC = cert;
```

The query defining the view involves two relations



Querying Views



- Query a view as if it were a base table.
 - Also: a limited ability to modify views if it makes sense as a modification of one underlying base table.

- Example query:

```
SELECT title
FROM ParamountMovies
WHERE year = 1979;
```




Querying Views: an example

ParamontMovie(movieTitle, movieYear)
Renaming Attributes

VIEW definitions

```
CREATE VIEW ParamountMovies AS
  SELECT title, year
  FROM Movie
  WHERE studioName = 'Paramount' ;
```

QUERY on VIEW

```
SELECT title
FROM ParamountMovies WHERE year = 1979;

Has the same effect as(equivalent query) on base table

SELECT title
FROM Movie
WHERE studioName = 'Paramount' AND year = 1979;
```



Querying Views: an example

SELECT DISTINCT starName
FROM ParamountMovies, StarsIn
WHERE title = movieTitle AND year = movieYear;

View

Base table

Queries involving both views and base tables

Equivalent query on table:

Interpreting the use of a virtual view as a subquery

```
SELECT DISTINCT starName
FROM (SELECT title, year
      FROM Movie
      WHERE studioName = 'Paramount'
      ) pm, StarsIn
WHERE pm.title = movieTitle AND pm.year = movieYear;
```



OUTLINES



8.1

Virtual Views

8.2

Modifying Views



Modifying Views

- In limited circumstances it is possible to execute an **Insertion**, **Deletion**, or **Update** to a view
- Modifying Data Through a View
 - **updatable views**, translate the modification of the view into an equivalent modification on a base table: and the modification can be done to the base table instead



Modifying Views

- View removal

DROP VIEW ParamountMovies;

DROP TABLE Movies;

Deletion the base table and make
all the views on it unusable



updatable views

- permit modifications on views that
 - Defined by selecting (using **SELECT**, not **SELECT DISTINCT**) some attributes from **one relation R**
 - The **WHERE** clause must not involve R in a subquery
 - The **FROM** clause can only consist of one occurrence of R and no other relation
 - (optional) The list of **SELECT** must includes enough attributes that for every tuple inserted into the view, we can fill the other attributes out with NULL values or the proper default
- Any modifications are translated into an equivalent modifications on a base table.



Example: Insertion

```
CREATE VIEW ParamountMovies AS  
  SELECT title, year  
  FROM Movie WHERE studioName = 'Paramount' ;
```

```
INSERT INTO ParamountMovies  
  VALUES ('Star Trek', 1979);
```

What happen? **Can not be seen in View!**

title	year	length	inColor	studioName	producerC#
Star Trek	1979	0	NULL	NULL	NULL



Example: Intersion

To fix this problem[includes enough attributes]

```
CREATE VIEW ParamountMovies AS  
    SELECT title, year, studioName  
    FROM Movie WHERE studioName = 'Paramount' ;
```

```
INSERT INTO ParamountMovie  
    VALUES('Star Trek', 1979, 'Paramount');
```

title	year	length	inColor	studioName	producerC#
Star Trek	1979	0	NULL	Paramount	NULL



Example: Deletion

Delete from the updatable ParamountMovie view all movies with "Trek" in their titles.

```
DELETE FROM ParamountMovies  
WHERE title LIKE '%Trek%'
```

This deletion is translated into an equivalent deletion on the Movie base table: the only difference is that **the condition defining the view ParamountMovies is added to the conditions of the WHERE clause.**

```
DELETE FROM Movies  
WHERE title LIKE '%Trek%'  
AND studioName = 'Paramount' ;
```



Example: Update

An update on an updatable view is passed through to the underlying relation.

```
UPDATE ParamountMovies
```

```
SET year = 1979
```

```
WHERE title = 'Star Trek the Movie';
```

is equivalent to the base-table update

```
UPDATE Movies
```

```
SET year = 1979
```

```
WHERE title = 'Star Trek the Movie'
```

```
AND studioName = 'Paramount' ;
```



Question

- Why some views are not updatable?



Triggers on Views

- Generally, it is impossible to modify a virtual view, because it doesn't exist.
- But an **INSTEAD OF** trigger lets us intercept view modifications in a way that makes sense.



Example

```
CREATE VIEW ParamountMovies AS
  SELECT title, year
  FROM Movies WHERE studioName = 'Paramount';
```

The view is updatable, but it has the unexpected flaw that when we insert a tuple into ParamountMovies.

A better solution is to create an instead-of trigger on this view.

```
CREATE TRIGGER ParamountInsert
  INSTEAD OF INSERT ON ParamountMovies
  REFERENCING
    NEW ROW AS NewRow
  FOR EACH ROW
  INSERT INTO Movies(title, year, stdioName)
  VALUES(NewRow.title, NewRow.year, 'Paramount');
```



OUTLINES



8.3

Indexes in SQL

8.4

Selection of Indexes

8.5

Materialized Views



OUTLINES



8.3

Indexes in SQL

8.4

Selection of Indexes

8.5

Materialized Views



Indexes

- *Index* = data structure used to speed access to tuples of a relation, given values of one or more attributes.
- Query (SELECT) more efficient
- Could be a hash table, but in a DBMS it is always a balanced search tree with giant nodes (a full disk page) called a *B-tree*.



Motivation

- When relations are very large,
 - It becomes expensive to scan all the tuples of a relation to find those (perhaps very few) tuples that match a given condition.

SELECT *

FROM Movies

WHERE studioName = 'Disney' AND year = 1990;



Motivation

- Indexes may also be useful in queries that involve a join,
- How do indexes on title of Movies and cert of MovieExec work?

SELECT name

FROM Movies, MovieExec

WHERE title = 'Star Wars' AND producerC = cert;



Declaring & Deleting Indexes

- Declaring

- No standard!

- Typical syntax:

```
CREATE INDEX YearIndex ON Movies(year);
```

```
CREATE INDEX KeyIndex ON Movies(title, year);
```

- Deleting

```
DROP INDEX YearIndex;
```

```
DROP INDEX KeyIndex;
```



Using Indexes

- Given a value **v**, the index takes us to only those tuples that have **v** in the attribute(s) of the index.
- **Example:**
 - SQL queries that **specify a year** may be executed by the SQL query processor in such a way that **only** those tuples of Movies **with the specified year** are ever examined
 - Decrease in the time needed to answer the query



Using Indexes

- Multiattribute index
 - The key is the concatenation of the attributes in some order
 - We can even use this index to find all the tuples with a given value in the first of the attributes
 - The choice of the order in which the attributes are listed

```
CREATE INDEX KeyIndex ON  
Movies(title, year);
```



OUTLINES



8.3

Indexes in SQL

8.4

Selection of Indexes

8.5

Materialized Views



Selection of Indexes



- Choosing which indexes to create is a **trade-off**
 - Index on an attribute may speed up greatly the execution of query
 - in which a value, or range of values, is specified for that attribute,
 - and may speed up joins involving that attribute as well
 - Index makes insertions, deletions and updates to that relation more complex and time-consuming.
 - modification(update, delete, insert) is about twice as expensive as selection
 - each modification of R forces to change any index on that modified attribute(s) of R.



A Simple Cost Model



- Time spent answering a query
 - To examine even one tuple requires that the whole page be brought into main memory
 - It costs little more time to examine all the tuples on a page than to examine only one
 - We assume every page we need must be retrieved from the disk



Some Useful Indexes

- The most useful index we can put on a relation is an index on its key
 - Queries in which a value for the key is specified are common. Thus, an index on the key will get used frequently
 - Since there is at most one tuple with a given key value, the index returns either nothing or one location for a tuple. Thus, at most one page must be retrieved



Some Useful Indexes

- Example

SELECT name

FROM Movies, MovieExec

WHERE title = 'Star Wars' AND producerC = cert;

- With the right indexes, the whole query might be done with as few as two page reads
- How to do this?



Some Useful Indexes

- When the index is not on a key
 - It may or may not be able to improve the time
- Two situations in which an index can be effective, even if it is not on a key
 - If the attribute is almost a key; that is, relatively few tuples have a given value for that attribute
 - If the tuples are “clustered” on that attribute. We cluster a relation on an attribute by grouping the tuples with a common value for that attribute



Some Useful Indexes

- When the index is not on a key
 - It may or may not be able to improve the time

- Example

SELECT name

FROM Movies, MovieExec

WHERE title = 'Star Wars' AND producerC = cert;

**CREATE INDEX TitleIndex ON
Movies(title) ;**

What happens compared with the index on the key?



Some Useful Indexes

- When the index is not on a key
 - It may or may not be able to improve the time

- Example

```
SELECT * FROM Movies WHERE year = 1990;
```

- The only index we have on Movies

```
CREATE INDEX YearIndex ON Movies(year);
```

- The tuples are **not clustered** by year
 - Say stored alphabetically by title
- The tuples are **clustered** on year



Calculating the Best Indexes

- One must be very careful how one estimates the relative frequency of modifications and queries
 - Modification, twice as expensive as query
 - The most common query and modification forms



Calculating the Best Indexes

- Depend on which queries and modifications are most likely to be performed on the Database.

- **Example** StarsIn(movieTitle, movieYear, starName)

Q1: **SELECT** movieTitle, movieYear

FROM StarsIn

WHERE starName = *s*;

Q2: **SELECT** starName

FROM StarsIn

WHERE movieTitle = *t* **AND** movieYear = *y*;

I1: **INSERT INTO** StarsIn **VALUES** (*t*, *y*, *s*);

Basic assumptions:

StarsIn is stored in 10 disk blocks (pages)

On the average, a star has appeared in 3 movies and a movie has 3 stars;



Example

- Costs associated with the three actions, as a function of which indexes are selected

Action	No Index	Star Index	Movie Index	Both Indexes
Q_1	10	4	10	4
Q_2	10	10	4	4
I	2	4	4	6
	$2 + 8p_1 + 8p_2$	$4 + 6p_2$	$4 + 6p_1$	$6 - 2p_1 - 2p_2$

the fraction of Q_1 is p_1 ;

the fraction of Q_2 is p_2 ,

the fraction of I is $1-p_1-p_2$.

- $p_1=p_2=0.1$, $(2+8p_1+8p_2)$ is the smallest
- $p_1=p_2=0.4$, $(6-2p_1-2p_2)$ is the smallest
- $p_1=0.5$, $p_2=0.1$, $(4+6p_2)$ is the smallest
- $p_1=0.1$, $p_2=0.5$, $(4+6p_1)$ is the smallest



Automatic Selection of Indexes to Create

- “Tuning” a database is a process
 - That includes not only index selection,
 - But the choice of many different parameters
- How the index-selection portion of tuning advisors work
 - To establish the query workload
 - The designer may be offered the opportunity to specify some constraints, e.g., indexes that must, or must not, be chosen
 - The tuning advisor generates a set of possible candidate indexes, and evaluates each one
 - The index set resulting in the lowest cost is created



Automatic Selection of Indexes to Create

- A “greedy” approach to choosing indexes
 - Initially, with no indexes selected, evaluate the benefit of each of the candidate indexes. If at least one provides positive benefit, then choose that index
 - Then, reevaluate the benefit of each of the remaining candidate indexes, assuming that the previously selected index is also available. Again, choose the index that provides the greatest benefit
 - In general, repeat the evaluation of candidate indexes under the assumption that all previously selected indexes are available. Pick the index with maximum benefit, until no more positive benefits can be obtained



OUTLINES



8.3

Indexes in SQL

8.4

Selection of Indexes

8.5

Materialized Views



Materialized Views

- If a view is used frequently enough, it may even be efficient to **materialize** it. (**More Efficient!!**)
 - To maintain its values at all times
 - Cost: recompute the **Materialized View** each time one of the underlying base tables changes.
 - All the changes to the Materialized View are **incremental**

```
CREATE MATERIALIZED VIEW MovieProd AS
SELECT title, year, name
FROM Movies, MovieExec
WHERE producerC = cert;
```



Materialized Views

All the changes to the Materialized View are *incremental*.
It is more efficient than a re-execution of the query.

Scene 1) Insert a new movie into Movies(title, year, producerC):
`INSERT INTO Movies VALUES("Kill Bill", 2003, 23456);`

Then we only need to look up cert=23456 in MovieExec (for cert is the key). So
`SELECT name FROM MovieExec WHERE cert = 23456;`

As this query return name = "Quentin Tarantino", then the DBMS can insert the proper tuple into MovieProd by

`INSERT INTO MovieProd VALUES('Kill Bill', 2003, 'Quentin Tarantino');`



Materialized Views

Scene 2) Delete a movie from Movies (title, year, producerC):

DELETE FROM Movies WHERE title = 'Dumb & Dumber' and year = 1994;

The DBMS will:

DELETE FROM MovieProd WHERE title = 'Dumb & Dumber' and year = 1994;

Scene 3) Update a movie in Movies (title, year, producerC):

UPDATE Movies SET title = 'Dumb & Dumber 2'

WHERE title='Dumb & Dumber' AND year = 1994;

The DBMS will:

UPDATE Moviesprod SET title = 'Dumb & Dumber 2'

WHERE title='Dumb & Dumber' AND year = 1994;

Scene 4) Insert a tuple into MovieExec (cert, name)

INSERT INTO MovieExec VALUES(34567, 'Max Bialystock');

The DBMS will:

INSERT INTO MovieProd

SELECT title, year, "Max Bialystock" FROM Movies WHERE producerC=34567;



Materialized Views

Scene 5) Delete a tuple from MovieExec

```
DELETE FROM MovieExec WHERE cert = 45678;
```

The DBMS will:

```
DELETE FROM MovieProd
```

```
WHERE (title, year) IN
```

```
(SELECT title, year FROM Movies WHERE producerC = 45678);
```

Scene 6) Update to MovieExec involving cert

```
UPDATE MovieExec SET cert = 12345 where cert = 34567
```

The DBMS will:

```
DELETE FROM MovieProd
```

```
WHERE (title, year) IN (
```

```
SELECT title, year FROM Movies, MovieExec
```

```
WHERE cert = 34567 AND cert = producerC);
```

```
INSERT INTO MovieProd
```

```
SELECT title, year, name
```

```
FROM Movies, MovieExec
```

```
WHERE cert = 12345 AND cert = producerC;
```

```
SELECT name  
FROM MovieExec  
WHERE cert = 45678
```

Can we delete/Update
all movies from MovieProd
that have the producer
name?



Periodic Maintenance

- **Problem:** each time a base table changes, the materialized view may change.
 - Cannot afford to recompute the view with each change.
- **Solution:** Periodic reconstruction of the materialized view, which is otherwise “out of date.” (typically each night)



Example: A Data Warehouse

- Wal-Mart stores every sale at every store in a database.
- Overnight, the sales for the day are used to update a *data warehouse* = materialized views of the sales.
- The warehouse is used by analysts to predict trends and move goods to where they are selling best.



Rewriting Queries to Use Materialized Views



- A materialized view V defined by

```
SELECT  $L_V$   
FROM  $R_V$   
WHERE  $C_V$ 
```

- Suppose we have a query Q

```
SELECT  $L_Q$   
FROM  $R_Q$   
WHERE  $C_Q$ 
```



Rewriting Queries to Use Materialized Views



Here are the conditions under which we can replace part of the query Q by the view V .

1. The relations in list R_V all appear in the list R_Q .
2. The condition C_Q is equivalent to C_V AND C for some condition C . As a special case, C_Q could be equivalent to C_V , in which case the “AND C ” is unnecessary.
3. If C is needed, then the attributes of relations on list R_V that C mentions are attributes on the list L_V .
4. Attributes on the list L_Q that come from relations on the list R_V are also on the list L_V .



Rewriting Queries to Use Materialized Views



If all these conditions are met, then we can rewrite Q to use V , as follows:

- a) Replace the list R_Q by V and the relations that are on list R_Q but not on R_V .
- b) Replace C_Q by C . If C is not needed (i.e., $C_V = C_Q$), then there is no **WHERE** clause.



Rewriting Queries to Use Materialized Views



- Example:

```
SELECT starName
FROM StarsIn, Movies, MovieExec
WHERE movieTitle= title AND movieYear = year AND producer = cert
      AND name= 'Max Bialystock'
```

- Rewrite

```
CREATE MATERIALIZED VIEW MovieProd AS
```

```
  SELECT title, name, year
     FROM Movies, MovieExec
     WHERE producer = cert
```

Movies, MovieExec

AND producer = cert

```
SELECT starName
   FROM StarsIn, MovieProd
  WHERE movieTitle= title AND movieYear = year AND name= 'Max Bialystock'
```



Automatic Creation of Materialized Views

- The ideas for indexes can apply as well to materialized views
 - To establish or approximate the query workload
 - An automated materialized-view-selection advisor needs to generate candidate views
 - Far more difficult than generating candidate indexes
 - All modern optimizers know how to exploit indexes, but not all can exploit materialized views