



LECTURE 8

CONCURRENCY CONTROL (PART 1/3)



NIKON D90 F3.5 3s ISO320



C5000Z F2.8 1/800s ISO50





Concurrency Control




- Motivation
 - Interactions among concurrently executing transactions
 - can cause the database state to become inconsistent
 - even when the transactions individually preserve correctness of the state, and there is no system failure
 - The timing of individual steps of different transactions needs to be regulated



Concurrency Control

- **Concurrency Control**
 - The process of assuring that transactions preserve consistency when executing simultaneously.
- **The scheduler**
 - takes read/write requests from transactions and either **executes them in buffers** or **delay them**





Concurrency Control

- Concurrency Control
 - How to assure that concurrently executing transactions preserve correctness of the database state
 - Serializability and conflict-serializability
 - Locking
 - Two-phase locking



OUTLINES



18.1 Serial and Serializable Schedules



18.2 Conflict-Serializability



OUTLINES



18.1

Serial and Serializable Schedules

18.2

Conflict-Serializability



Recall: Correctness Principle

- The *Correctness Principle*: if a transaction executes in the absence of any other transactions or system errors, and it starts with the DB in a consistent state, then the DB is also in a consistent state when the transaction ends.
- In practice
 - Transactions often run concurrently with other transactions, so the correctness principle does not apply directly



Schedules

- Schedule

- A sequence of the important actions taken by one or more transactions

T_1	T_2
READ(A,t)	READ(A,s)
t := t+100	s := s*2
WRITE(A,t)	WRITE(A,s)
READ(B,t)	READ(B,s)
t := t+100	s := s*2
WRITE(B,t)	WRITE(B,s)



Serial Schedules

- A schedule is **serial**
 - If its actions consist of all the actions of one transaction, then all the actions of another transaction, and so on.
 - No mixing of the actions is allowed.



Serial Schedules

- A schedule is **serial**
 - T_1 precedes T_2

T_1	T_2	A	B
		25	25
READ(A,t)			
t := t+100			
WRITE(A,t)		125	
READ(B,t)			
t := t+100			
WRITE(B,t)			125
	READ(A,s)		
	s := s*2		
	WRITE(A,s)	250	
	READ(B,s)		
	s := s*2		
	WRITE(B,s)		250



Serial Schedules

- A schedule is **serial**

- T_2 precedes T_1

- In general, we would not expect the final state of a database to be independent of the order of transactions

T_1	T_2	A	B
		25	25
	READ(A,s)		
	s := s*2		
	WRITE(A,s)	50	
	READ(B,s)		
	s := s*2		
	WRITE(B,s)		50
READ(A,t)			
t := t+100			
WRITE(A,t)		150	
READ(B,t)			
t := t+100			
WRITE(B,t)			150



Serializable Schedules

- The correctness principle tells us
 - Every serial schedule will preserve consistency of the database state.
 - Are there any other schedules that also are guaranteed to preserve consistency?
- Serializable schedule
 - A schedule S is *serializable* if there is a serial schedule S' such that for every initial database state, the effects of S and S' are the same



Serializable Schedules

- Serializable schedule

- Example:

Which is the same as for
the serial schedule

(T_1, T_2)

T_1	T_2	A	B
		25	25
READ(A,t) t := t+100 WRITE(A,t)		125	
	READ(A,s) s := s*2 WRITE(A,s)	250	
READ(B,t) t := t+100 WRITE(B,t)			125
	READ(B,s) s := s*2 WRITE(B,s)		250



Serializable Schedules

- Serializable schedule

- Example

which is not serializable

- This schedule is the sort of behavior that concurrency control mechanisms must avoid

T_1	T_2	A	B
		25	25
READ(A,t) t := t+100 WRITE(A,t)		125	
	READ(A,s) s := s*2 WRITE(A,s) READ(B,s) s := s*2 WRITE(B,s)	250	
			50
READ(B,t) t := t+100 WRITE(B,t)			150



The Effect of Transaction Semantics

- Example

- Coincidentally, it also results from the serial schedule (T_2, T_1)

- Assume:

- Any database element A that a transaction T writes is given a value that depends on the database state in such a way that no arithmetic coincidences occur.

T_1	T_2	A	B
		25	25
READ(A, t) $t := t+100$ WRITE(A, t)		125	
	READ(A, s) $s := s+200$ WRITE(A, s) READ(B, s) $s := s+200$ WRITE(B, s)	325	
			225
READ(B, t) $t := t+100$ WRITE(B, t)			325



A Notation for Transactions and Schedules

- If we assume “no coincidences”
 - Then only the reads and writes performed by the transactions matter,
 - not the actual values involved
- Notation

$r_T(X)$ transaction T reads database element X

$w_T(X)$ transaction T writes database element X

$r_i(X)$ and $w_i(X)$ synonyms for $r_{Ti}(X)$ and $w_{Ti}(X)$



A Notation for Transactions and Schedules

- Example

T_1	T_2
READ(A,t)	READ(A,s)
t := t+100	s := s*2
WRITE(A,t)	WRITE(A,s)
READ(B,t)	READ(B,s)
t := t+100	s := s*2
WRITE(B,t)	WRITE(B,s)

$T_1: r_1(A); w_1(A); r_1(B); w_1(B);$

$T_2: r_2(A); w_2(A); r_2(B); w_2(B);$



A Notation for Transactions and Schedules

- Example

T_1	T_2	A	B
		25	25
READ(A,t)			
t := t+100			
WRITE(A,t)		125	
	READ(A,s)		
	s := s*2		
	WRITE(A,s)	250	
READ(B,t)			
t := t+100			
WRITE(B,t)			125
	READ(B,s)		
	s := s*2		
	WRITE(B,s)		250

$r_1(A); w_1(A); r_2(A); w_2(A); r_1(B); w_1(B); r_2(B); w_2(B);$



A Notation for Transactions and Schedules

- An schedule S
 - of a set of transactions \mathcal{T} is a sequence of actions,
 - in which for each transaction T_i in \mathcal{T} , the actions of T_i appear in S in the same order that they appear in the definition of T_i itself



OUTLINES



18.1

Serial and Serializable Schedules

18.2

Conflict-Serializability



Conflicts

- **Conflict**
 - A pair of **consecutive** actions in a schedule such that,
 - If their order is **interchanged**, then behavior of at least one of the transactions involved can change



Conflicts

- Most pairs of actions do *not* conflict
 1. $r_i(X); r_j(Y)$ is never a conflict, even if $X=Y$
 2. $r_i(X); w_j(Y)$ is not a conflict if $X \neq Y$
 3. $w_i(X); r_j(Y)$ is not a conflict if $X \neq Y$
 4. $w_i(X); w_j(Y)$ is not a conflict if $X \neq Y$



Conflicts

- May *not* swap the order of actions
 1. Two actions of the **same** transaction:
 $r_i(X); r_i(Y)$, $r_i(X); w_i(Y)$, $w_i(X); r_i(Y)$, $w_i(X); w_i(Y)$
 2. A read and a write of the **same** database element by **different** transactions
 $r_i(X); w_j(X)$, $w_i(X); r_j(X)$
 3. Two writes of the **same** database element by **different** transactions
 $w_i(X); w_j(X)$



Conflicts



- Any two actions of different transactions may be swapped unless
 1. They involve the same database element, and
 2. At least one is a write



Conflicts

- **Conflict-equivalent**
 - Two schedules are **conflict-equivalent** if they can be turned one into the other **by a sequence of nonconflicting swaps of adjacent actions**
- **Conflict-serializable**
 - A schedule is **conflict-serializable** if it is **conflict-equivalent** to a **serial schedule**



Conflicts

- Conflict-serializable
 - Conflict-serializability is a sufficient condition for serializability
 - A conflict-serializable schedule is a serializable schedule



Conflicts

- Example

$r_1(A); w_1(A); r_2(A); w_2(A); r_1(B); w_1(B); r_2(B); w_2(B);$

$r_1(A); w_1(A); r_2(A); \underline{w_2(A)}; \underline{r_1(B)}; w_1(B); r_2(B); w_2(B);$
 $r_1(A); w_1(A); \underline{r_2(A)}; \underline{r_1(B)}; w_2(A); w_1(B); r_2(B); w_2(B);$
 $r_1(A); w_1(A); r_1(B); \underline{r_2(A)}; \underline{w_2(A)}; \underline{w_1(B)}; r_2(B); w_2(B);$
 $r_1(A); w_1(A); r_1(B); \underline{r_2(A)}; \underline{w_1(B)}; w_2(A); r_2(B); w_2(B);$
 $r_1(A); w_1(A); r_1(B); w_1(B); \underline{r_2(A)}; w_2(A); r_2(B); w_2(B);$



Precedence Graphs

- To decide whether or not it is conflict-serializable
- T_1 takes precedence of T_2
 - Given a schedule S , involving transactions T_1 and T_2
 - $T_1 <_S T_2$ if there are actions A_1 of T_1 and A_2 of T_2 , such that
 - A_1 is ahead of A_2 in S ,
 - Both A_1 and A_2 involve the same database elements
 - At least one of A_1 and A_2 is a write action



Precedence Graphs

- Precedence Graphs

- Nodes

- The transactions of a schedule S
- Label the node for T_i by only integer i

- Edges


- There is an arc from node i to node j if $T_i <_S T_j$



Precedence Graphs

- Example

$S: r_2(A); r_1(B); w_2(A); r_3(A); w_1(B); w_3(A); r_2(B); w_2(B);$



The precedence graph for the schedule S



Test for Conflict-Serializability

- To tell whether a schedule S is conflict-serializable
 1. Construct the precedence graph for S
 2. Ask if there are any cycles
 - If so, then S is not conflict-serializable
 - If not, then S is conflict-serializable
- Topological order of the nodes is a conflict-equivalent serial order



Conflicts

- **Conflict-serializability** is not necessary for serializability

– Counterexample:

$S_1: w_1(Y); w_1(X); w_2(Y); w_2(X); w_3(X);$ serial

$S_2: w_1(Y); w_2(Y); w_2(X); w_1(X); w_3(X);$

S_2 is serializable

S_2 is not conflict -serializable



Test for Conflict-Serializability

- Example

$S: r_2(A); r_1(B); w_2(A); r_3(A); w_1(B); w_3(A); r_2(B); w_2(B);$



The precedence graph for the schedule S

The precedence graph is **acyclic**, so the schedule S is **conflict-serializable**


How to swap? $S': r_1(B); w_1(B); r_2(A); w_2(A); r_2(B); w_2(B); r_3(A); w_3(A);$



Test for Conflict-Serializability

- Example

$S_1: r_2(A); r_1(B); w_2(A); r_2(B); r_3(A); w_1(B); w_3(A); w_2(B);$



The precedence graph is **cyclic**, so the schedule S_1 is not **conflict-serializable**



Why the Precedence-Graph Test Works

• Why the Precedence-Graph Test Works?

Proof

1. If there is a cycle involving n transactions
 $T_1 \rightarrow T_2 \rightarrow \dots \rightarrow T_n \rightarrow T_1$
2. Then in the hypothetical serial order, the actions of T_1 must precede those of T_2 , and so on, up to T_n
3. But the actions of T_n are also required to precede those of T_1
 - If there is a cycle in the precedence graph, then the schedule is not conflict-serializable



Why the Precedence-Graph Test Works

- Why the Precedence-Graph Test Works?

Proof: If the precedence graph has no cycles, then the schedule is conflict-serializable

BASIS:

- If $n=1$, i.e., there is only one transaction in the schedule, then the schedule is already serial, and therefore surely conflict-serializable



Why the Precedence-Graph Test Works

- Why the Precedence-Graph Test Works?

Proof: If the precedence graph has no cycles, then the schedule is conflict-serializable

INDUCTION:

- Let the schedule S consist of the actions of n transactions T_1, T_2, \dots, T_n
- We suppose that S has an acyclic precedence graph.
 - If a finite graph is acyclic, then there is at least one node that has no arcs in;
 - Let the node i corresponding to transaction T_i be such a node



Why the Precedence-Graph Test Works

- **Why the Precedence-Graph Test Works?**

Proof: If the precedence graph has no cycles, then the schedule is conflict-serializable

INDUCTION:

- Since there are no arcs into node i , there can be no action A in S that
 1. Involves any transaction T_j other than T_i ,
 2. Precedes some action of T_i , and
 3. Conflicts with that action
- For if there were, we should have put an arc from node j to node i in the precedence graph



Why the Precedence-Graph Test Works

- Why the Precedence-Graph Test Works?

Proof: If the precedence graph has no cycles, then the schedule is conflict-serializable

INDUCTION:

- It is thus possible to swap all the actions of T_i , keeping them in order, but moving them to the front of S
- The schedule has now taken the form
(Actions of T_i) (Actions of the other $n-1$ transactions)



Why the Precedence-Graph Test Works

- **Why the Precedence-Graph Test Works?**

Proof: If the precedence graph has no cycles, then the schedule is conflict-serializable

INDUCTION:

- Let us now consider the tail of S – the actions of all transactions other than T_i
- Since these actions maintain the same relative order that they did in S
 - The precedence graph for the tail is the same as the precedence graph for S , except that the node for T_i and any arcs out of that node is missing



Why the Precedence-Graph Test Works

- **Why the Precedence-Graph Test Works?**

Proof: If the precedence graph has no cycles, then the schedule is conflict-serializable

INDUCTION:

- Since the original precedence graph was acyclic, and deleting nodes and arcs cannot make it cyclic, we conclude that the tail's precedence graph is acyclic
- Moreover, since the tail involves $n-1$ transactions, the inductive hypothesis applies to it



Why the Precedence-Graph Test Works

- **Why the Precedence-Graph Test Works?**

Proof: If the precedence graph has no cycles, then the schedule is conflict-serializable

INDUCTION:

- Thus, we know we can reorder the actions of the tail using legal swaps of adjacent actions to turn it into a serial schedule.
- Now, S itself has been turned into a serial schedule, with the actions of T_i first and the actions of the other transactions following in some serial order



Why the Precedence-Graph Test Works

- Why the Precedence-Graph Test Works?

Proof: If the precedence graph has no cycles, then the schedule is conflict-serializable

INDUCTION:

- The induction is complete, and we conclude that every schedule with an acyclic precedence graph is conflict-serializable.