



LECTURE 08

CONCURRENCY CONTROL (PART 2/3)



NIKON D90 F3.5 3s ISO320



C5000Z F2.8 1/800s ISO50





OUTLINES



18.3 Enforcing Serializability by Locks



Introduction


- The most common architecture for a scheduler
 - “locks” are maintained on database elements to prevent unserializable behavior
- In this section
 - The concept of locking
 - One kind of lock
 - Which transactions must obtain on a database element if they want to perform any operation on that element



Locks

- The responsibility of the scheduler
 - Take requests from transactions
 - Either allow them to operate on the database
 - Or block the transaction until such time as it is safe to allow it to continue

A lock table
will be used
to guide this
decision





Locking Scheduler

- A locking scheduler
 - Enforces conflict-serializability
 - Which is more stringent condition than serializability
 - Transactions must **request** and **release** locks
 - The structure of transactions
 - The structure of schedules



Locking Scheduler

- A locking scheduler
 - Consistency of Transactions:
 - Actions and locks must relate in the expected ways
 - A transaction can only read or write an element if it previously was granted a lock on that element and has not yet released the lock
 - If a transaction locks an element, it must later unlock that element
 - Legality of Schedules:
 - Locks must have their intended meaning:
 - No two transactions may have locked the same element without one having first released the lock



Notations

- Notations for locking and unlocking

$l_i(X)$

- Transaction T_i requests a lock on database element X

$u_i(X)$

- Transaction T_i releases (unlocks) its lock on database element X



Consistency Condition & Legality of Schedules

- Consistency Condition

- Whenever a transaction T_i has an action $r_i(X)$ or $w_i(X)$, then there is a previous action $l_i(X)$ with no intervening action $u_i(X)$, and there is a subsequent $u_i(X)$

- Legality of Schedules

- If there are actions $l_i(X)$ followed by $l_j(X)$ in a schedule, then somewhere between these actions there must be an action $u_i(X)$



Example

- Each of these transactions is consistent

$T_1: l_1(A); r_1(A); A := A+100; w_1(A); u_1(A); l_1(B); r_1(B); B := B+100; w_1(B); u_1(B);$

$T_2: l_2(A); r_2(A); A := A*2; w_2(A); u_2(A); l_2(B); r_2(B); B := B*2; w_2(B); u_2(B);$



Example

- A legal schedule of consistent transactions
- But it is not serializable

- The additional condition
- “Two-phase locking”
 - That we need to assure that legal schedules are conflict-serializable

T_1	T_2	A	B
		25	25
$l_1(A); r_1(A);$ $A := A+100;$ $w_1(A); u_1(A);$		125	
	$l_2(A); r_2(A);$ $A := A*2;$ $w_2(A); u_2(A);$ $l_2(B); r_2(B);$ $B := B*2;$ $w_2(B); u_2(B);$	250	
			50
$l_1(B); r_1(B);$ $B := B+100;$ $w_1(B); u_1(B);$			150



The Locking Scheduler

- The Locking Scheduler
 - It is the job of a scheduler based on locking to grant requests **if and only if** the request will result in a legal schedule
 - If a request is not granted, the requesting transaction is delayed
 - It waits until the scheduler grants its request at a later time



The Locking Scheduler

- Lock Table

- The scheduler has a **lock table** that tells, for every database element, the transaction (if any) that currently holds a lock on that element
- Only one kind of lock: simple case
 - A relation **Locks (element, transaction)**
 - Consisting of pairs (X, T) such that transaction T currently has a lock on database element X



Example

- Sometimes it is not possible to grant requests

$T_1: l_1(A); r_1(A); A := A+100; w_1(A); l_1(B); u_1(A); r_1(B); B := B+100; w_1(B); u_1(B);$

$T_2: l_2(A); r_2(A); A := A*2; w_2(A); l_2(B); u_2(A); r_2(B); B := B*2; w_2(B); u_2(B);$

T_1	T_2	A	B
		25	25
$l_1(A); r_1(A);$ $A := A+100;$ $w_1(A); l_1(B); u_1(A);$		125	
	$l_2(A); r_2(A);$ $A := A*2;$ $w_2(A);$ $l_2(B)$ Denied	250	
$r_1(B); B := B+100;$ $w_1(B); u_1(B);$			125
	$l_2(B); u_2(A); r_2(B);$ $B := B*2;$ $w_2(B); u_2(B);$		250



Two-Phase Locking

- Two-Phase Locking (2PL)
 - Guarantee that a legal schedule of consistent transactions is conflict-serializable
- 2PL
 - In every transaction, all lock actions precede all unlock actions



Two-Phase Locking

- 2PL
 - In every transaction, all lock actions precede all unlock actions
- “Two-Phase”
 - The first phase: locks are obtained
 - The second phase: locks are relinquished



Two-Phase Locking

- 2PL
 - is a condition on the order of actions in a transaction
 - A transaction that obeys the 2PL condition is said to be a two-phase-locked transaction, or 2PL transaction



Example

- The transactions do not obey the two-phase locking rule

T_1	T_2	A	B
		25	25
$l_1(A); r_1(A);$ $A := A+100;$ $w_1(A); u_1(A);$		125	
	$l_2(A); r_2(A);$ $A := A*2;$ $w_2(A); u_2(A);$	250	
	$l_2(B); r_2(B);$ $B := B*2;$ $w_2(B); u_2(B);$		50
$l_1(B); r_1(B);$ $B := B+100;$ $w_1(B); u_1(B);$			150



Example

- The transactions **do** obey the **2PL** condition


T_1	T_2	A	B
		25	25
$l_1(A); r_1(A);$ $A := A+100;$ $w_1(A); l_1(B); u_1(A);$		125	
	$l_2(A); r_2(A);$ $A := A*2;$ $w_2(A);$ $l_2(B)$ Denied	250	
$r_1(B); B := B+100;$ $w_1(B); u_1(B);$			125
	$l_2(B); u_2(A); r_2(B);$ $B := B*2;$ $w_2(B); u_2(B);$		250



Why Two-Phase Locking Works

- Intuitively

- Each two-phase-locked transaction may be thought to execute **in its entirety** at the instant it issues its **first unlock** request
- There is always at least one **conflict-equivalent serial schedule** for a schedule ***S*** of **2PL transactions**
 - The one in which the transaction appear **in the same order** as their **first unlocks**





Why Two-Phase Locking Works

- How to convert any legal schedule S of consistent, two-phase-locked transactions to a conflict-equivalent serial schedule
 - An induction on the number of transactions in S
 - Issue of conflict-equivalence refers to the read and write actions only
 - As we swap the order of reads and writes, we ignore the lock and unlock actions



Why Two-Phase Locking Works

- An induction on the number of transactions in S
 - BASIS:
 - If $n = 1$, there is nothing to do; S is already a serial schedule



Why Two-Phase Locking Works

- An induction on the number of transactions in S

– INDUCTION:

- Suppose S involves n transactions T_1, T_2, \dots, T_n , and let T_i be the transaction with the first unlock action in the entire schedule S , say $u_i(X)$
- We claim it is possible to move all the read and write actions of T_i forward to the beginning of the schedule without passing any conflicting reads and writes



Why Two-Phase Locking Works

- An induction on the number of transactions in S

– INDUCTION:

- Consider some action of T_i , say $w_i(Y)$.
- Could it be preceded in S by some conflicting action, say $w_j(Y)$?
- If so, then in schedule S , actions $u_j(Y)$ and $l_i(Y)$ must intervene, in a sequence of actions
... $w_j(Y)$; ... ; $u_j(Y)$; ... ; $l_i(Y)$; ... ; $w_i(Y)$; ...



Why Two-Phase Locking Works

– INDUCTION:

- If so, then in schedule **S**, actions $u_j(Y)$ and $l_i(Y)$ must **intervene**, in a sequence of actions

... $w_j(Y)$; ... ; $u_j(Y)$; ... ; $l_i(Y)$; ... ; $w_i(Y)$; ...

- Since T_i is the first to unlock, $u_i(X)$ precedes $u_j(Y)$ in S; that is, S might look like:

... $w_j(Y)$; ... ; $u_i(X)$; ... ; $u_j(Y)$; ... ; $l_i(Y)$; ... ; $w_i(Y)$; ...

or $u_i(X)$ could even appear before $w_j(Y)$.

In any case, $u_i(X)$ appears before $l_i(Y)$,

which means that T_i is **not two-phase-locked**, as we assumed



Why Two-Phase Locking Works

– INDUCTION:

- While we have only argued the nonexistence of conflicting pairs of writes, the same argument applies to any pair of potentially conflicting actions, one from T_i and the other from another T_j
- We **conclude** that it is indeed possible to move all the actions of T_i forward to the beginning of S , using swaps of nonconflicting read and write actions, followed by restoration of the lock and unlock actions of T_i . That is, S can be written in the form

(Actions of T_i) (Actions of the other $n-1$ transactions)



Why Two-Phase Locking Works

– INDUCTION:

(Actions of T_i) (Actions of the other $n-1$ transactions)

- The tail of $n-1$ transactions is still a legal schedule of consistent, 2PL transactions,
- so the inductive hypothesis applies to it.
- We convert the tail to a conflict-equivalent serial schedule, and now all of S has been shown conflict-serializable



A Risk of Deadlock

- The potential for **deadlocks**
 - that is not solved by two-phase locking
 - where several transactions are forced by the scheduler to wait **forever** for a lock held by another transaction



A Risk of Deadlock

- Example

T_1	T_2	A	B
		25	25
$l_1(A); r_1(A);$	$l_2(B); r_2(B);$		
$A := A+100;$	$B := B*2;$		
$w_1(A);$	$w_2(B);$	125	
$l_1(B)$ Denied	$l_2(A)$ Denied		50